

Step 5: Performance Evaluation

- Load Test Data
 - Load the original dataset
 - Apply the same preprocessing pipeline used in training.ipynb
 - Extract the 20% test set (the same as used during training)
 - Load Trained Models & Pipeline
- Load the saved classification model
- Load the saved regression model
- Load the preprocessing pipeline

Evaluate Classification Model

- Generate predictions on the test set. Compute:
 - Accuracy
 - Precision, Recall, F1-score
 - Confusion Matrix
 - ROC-AUC Curve

Evaluate Regression Model

- Generate predictions on the test set. Compute:
 - R^2 (Coefficient of Determination)
 - MAE (Mean Absolute Error)
 - MSE (Mean Squared Error)
 - RMSE (Root Mean Squared Error)

Task 6. After completing all steps above, provide the following:

- Compare models and justify which one is better for each task.
- At least one visualizations per classification tasks (e.g., confusion matrix, ROC curve, precision-recall curves).

```
In [1]: import pandas as pd
import joblib
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score, recall_score
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_e
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
from utility import DataCleaner

df = pd.read_csv('Spotify_Song_Attributes.csv')
```

```
In [2]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10080 entries, 0 to 10079
Data columns (total 22 columns):
#   Column                Non-Null Count  Dtype
---  -
0   trackName             10080 non-null  object
1   artistName            10080 non-null  object
2   msPlayed              10080 non-null  int64
3   genre                 8580 non-null   object
4   danceability          9530 non-null   float64
5   energy                9530 non-null   float64
6   key                   9530 non-null   float64
7   loudness              9530 non-null   float64
8   mode                  9530 non-null   float64
9   speechiness           9530 non-null   float64
10  acousticness           9530 non-null   float64
11  instrumentalness       9530 non-null   float64
12  liveness               9530 non-null   float64
13  valence                9530 non-null   float64
14  tempo                  9530 non-null   float64
15  type                   9530 non-null   object
16  id                     9530 non-null   object
17  uri                    9530 non-null   object
18  track_href             9530 non-null   object
19  analysis_url           9530 non-null   object
20  duration_ms            9530 non-null   float64
21  time_signature         9530 non-null   float64
dtypes: float64(13), int64(1), object(8)
memory usage: 1.7+ MB
```

```
In [3]: data_cleaning_pipeline=joblib.load('cleaning_pipeline.pkl')
```

```
In [4]: df = data_cleaning_pipeline.transform(df)
```

```
/Users/chaotzuchieh/Documents/GitHub/project-1-TzuChieh_Chao/venv/lib/python3.12/site-packages/sklearn/pipeline.py:62: FutureWarning: This Pipeline instance is not fitted yet. Call 'fit' with appropriate arguments before using other methods such as transform, predict, etc. This will raise an error in 1.8 instead of the current warning.
  warnings.warn(
```

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 4261 entries, 1 to 5039
Data columns (total 22 columns):
#   Column                Non-Null Count  Dtype
---  -
0   trackName              4261 non-null   object
1   artistName             4261 non-null   object
2   msPlayed               4261 non-null   int64
3   genre                  4261 non-null   object
4   danceability           4261 non-null   float64
5   energy                 4261 non-null   float64
6   key                    4261 non-null   float64
7   loudness               4261 non-null   float64
8   mode                   4261 non-null   float64
9   speechiness            4261 non-null   float64
10  acousticness           4261 non-null   float64
11  instrumentalness        4261 non-null   float64
12  liveness               4261 non-null   float64
13  valence                4261 non-null   float64
14  tempo                  4261 non-null   float64
15  type                   4261 non-null   object
16  id                     4261 non-null   object
17  uri                    4261 non-null   object
18  track_href             4261 non-null   object
19  analysis_url           4261 non-null   object
20  duration_ms            4261 non-null   float64
21  time_signature         4261 non-null   float64
dtypes: float64(13), int64(1), object(8)
memory usage: 765.6+ KB
```

In [6]: `df.head()`

Out [6]:

	trackName	artistName	msPlayed	genre	danceability	energy	key
1	"In The Hall Of The Mountain King" from Peer G...	London Symphony Orchestra	1806234	british orchestra	0.475	0.130	7
2	#BrooklynBloodPop!	SyKo	145610	glitchcore	0.691	0.814	1
3	\$10	Good Morning	25058	experimental pop	0.624	0.596	4
4	(I Just) Died In Your Arms	Cutting Crew	5504949	album rock	0.625	0.726	11
5	(L)only Child	salem ilese	2237969	alt z	0.645	0.611	8

5 rows x 22 columns

In [7]: `classification_preprocessing_pipeline= joblib.load('cls_preprocessor.pkl')`
`regression_preprocessing_pipeline= joblib.load('reg_preprocessor.pkl')`

In [8]: `logreg_model = joblib.load('best_logreg_model_danceability.pkl')`
`rf_classifier = joblib.load('best_rf_model_danceability.pkl')`

```
linear_regression_model = joblib.load('linear_regression_model.pkl')
ridge_linear_model = joblib.load('ridge_model.pkl')
lasso_linear_model = joblib.load('lasso_model.pkl')
GridSearch_decision_tree = joblib.load('best_GridSearch_decision_tree_mod
RandomizedSearch_decision_tree = joblib.load('best_RandomizedSearch_decis
```

Classification

```
In [9]: cls_x_test = pd.read_csv('df_cls_x_test.csv')
        cls_y_test = pd.read_csv('df_cls_y_test.csv')
```


Logistic regression

```
In [10]: y1_pred_log_reg = logreg_model.predict(cls_x_test)
        y1_pred_log_reg_proba = logreg_model.predict_proba(cls_x_test)

        print('Logistic Regression')
        print('Accuracy:', accuracy_score(cls_y_test, y1_pred_log_reg))
        print('Precision:', precision_recall_fscore_support(cls_y_test, y1_pred_log_reg)[0])
        print('Recall:', precision_recall_fscore_support(cls_y_test, y1_pred_log_reg)[1])
        print('F1 Score:', precision_recall_fscore_support(cls_y_test, y1_pred_log_reg)[2])
        print('Confusion Matrix:', confusion_matrix(cls_y_test, y1_pred_log_reg))

        cm = confusion_matrix(cls_y_test, y1_pred_log_reg)
        sns.heatmap(cm, annot=True, fmt='d')
        plt.xlabel('Predicted')
        plt.ylabel('Actual')
        plt.title('Confusion Matrix')
        plt.show()
```

```
Logistic Regression
Accuracy: 0.7456037514654161
Precision: 0.5447154471544715
Recall: 0.29385964912280704
F1 Score: 0.3817663817663818
Confusion Matrix: [[569  56]
 [161  67]]
```




Random Forest Classifier

```
In [11]: y1_pred_rand_for=rf_classifier.predict(cls_x_test)
y1_pred_rand_for_proba=rf_classifier.predict_proba(cls_x_test)[:,-1]

print('Random Forest')
print('Accuracy:',accuracy_score(cls_y_test,y1_pred_rand_for))
print('Precision:',precision_recall_fscore_support(cls_y_test,y1_pred_rand_for)[0])
print('Recall:',precision_recall_fscore_support(cls_y_test,y1_pred_rand_for)[1])
print('F1 Score:',precision_recall_fscore_support(cls_y_test,y1_pred_rand_for)[2])
print('Confusion Matrix:',confusion_matrix(cls_y_test,y1_pred_rand_for))
cm=confusion_matrix(cls_y_test,y1_pred_rand_for)
sns.heatmap(cm,annot=True,fmt='d')
plt.xlabel('predicted')
plt.ylabel('actual')
plt.title('confusion matrix')
```

```
Random Forest
Accuracy: 0.7549824150058617
Precision: 0.5714285714285714
Recall: 0.3333333333333333
F1 Score: 0.42105263157894735
Confusion Matrix: [[568  57]
 [152  76]]
```

```
Out[11]: Text(0.5, 1.0, 'confusion matrix')
```



Regression

```
In [12]: reg_x_test = pd.read_csv('df_reg_x_test.csv')
reg_y_test = pd.read_csv('df_reg_y_test.csv')
```

Linear Regression

```
In [13]: y2_pred_lig_reg=linear_regression_model.predict(reg_x_test)
print('Linear Regression')
print('R2 Score:',r2_score(reg_y_test,y2_pred_lig_reg))
print('Mean Absolute Error:',mean_absolute_error(reg_y_test,y2_pred_lig_reg))
print('Mean Squared Error:',mean_squared_error(reg_y_test,y2_pred_lig_reg))
print('Root Mean Squared Error:',np.sqrt(mean_squared_error(reg_y_test,y2_pred_lig_reg)))
```

Linear Regression
 R2 Score: 0.3602056987416682
 Mean Absolute Error: 0.15229598307314432
 Mean Squared Error: 0.03626068026938245
 Root Mean Squared Error: 0.190422373342479

Linear Regression(Ridge)

```
In [14]: y2_pred_lig_reg=ridge_linear_model.predict(reg_x_test)
print('Linear Regression(Ridge)')
print('R2 Score:',r2_score(reg_y_test,y2_pred_lig_reg))
```

```
print('Mean Absolute Error:',mean_absolute_error(reg_y_test,y2_pred_lig_r
print('Mean Squared Error:',mean_squared_error(reg_y_test,y2_pred_lig_reg
print('Root Mean Squared Error:',np.sqrt(mean_squared_error(reg_y_test,y2
```

Linear Regression(Ridge)
R2 Score: 0.35670687196360706
Mean Absolute Error: 0.15419020882149784
Mean Squared Error: 0.03645897812678397
Root Mean Squared Error: 0.19094234241462518

Linear Regression(Lasso)

```
In [15]: y2_pred_lig_reg=lasso_linear_model.predict(reg_x_test)
print('Linear Regression(Lasso)')
print('R2 Score:',r2_score(reg_y_test,y2_pred_lig_reg))
print('Mean Absolute Error:',mean_absolute_error(reg_y_test,y2_pred_lig_r
print('Mean Squared Error:',mean_squared_error(reg_y_test,y2_pred_lig_reg
print('Root Mean Squared Error:',np.sqrt(mean_squared_error(reg_y_test,y2
```

Linear Regression(Lasso)
R2 Score: 0.35649835874578195
Mean Absolute Error: 0.15423930116020515
Mean Squared Error: 0.03647079572364069
Root Mean Squared Error: 0.19097328536641112

DecisionTreeRegressor(GridSearchCV)

```
In [16]: y2_pred_decision_tree=GridSearch_decision_tree.predict(reg_x_test)
print('Decision Tree')
print('R2 Score:',r2_score(reg_y_test,y2_pred_decision_tree))
print('Mean Absolute Error:',mean_absolute_error(reg_y_test,y2_pred_decis
print('Mean Squared Error:',mean_squared_error(reg_y_test,y2_pred_decisio
print('Root Mean Squared Error:',np.sqrt(mean_squared_error(reg_y_test,y2
```

Decision Tree
R2 Score: 0.2950428441908042
Mean Absolute Error: 0.16042207583291546
Mean Squared Error: 0.03995381950126051
Root Mean Squared Error: 0.19988451541142577

DecisionTreeRegressor(RandomizedSearchCV)

```
In [17]: y2_pred_decision_tree=RandomizedSearch_decision_tree.predict(reg_x_test)
print('Decision Tree')
print('R2 Score:',r2_score(reg_y_test,y2_pred_decision_tree))
print('Mean Absolute Error:',mean_absolute_error(reg_y_test,y2_pred_decis
print('Mean Squared Error:',mean_squared_error(reg_y_test,y2_pred_decisio
print('Root Mean Squared Error:',np.sqrt(mean_squared_error(reg_y_test,y2
```

Decision Tree
R2 Score: 0.2509975738287197
Mean Absolute Error: 0.16550395583307478
Mean Squared Error: 0.04245010848482428
Root Mean Squared Error: 0.20603424104945342

Analysis of Models

Classification

The results indicated that the Random Forest classifier outperformed the Logistic Regression model.

Regression

I implemented several machine learning regression models aimed at predicting valence to determine which model exhibited superior performance. However, it is noteworthy that none of the models yielded satisfactory results, with linear regression emerging as the most effective option.